

```
In [72]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from collections import Counter
```

Decision Tree Classifier

```
In [64]: class Node:
    def __init__(self, predicted_class):
        self.predicted_class = predicted_class
        self.feature_index = 0
        self.threshold = 0
        self.left = None
        self.right = None

class DecisionTreeClassifier:
    def __init__(self, max_depth=None):
        self.max_depth = max_depth

    def fit(self, X, y):
        self.n_classes_ = len(set(y))
        self.n_features_ = X.shape[1]
        self.tree_ = self._grow_tree(X, y)

    def predict(self, X):
        if np.isscalar(X):
            X = np.array([X])
        return [self._predict(inputs) for inputs in X]

    def _best_split(self, X, y):
        m = y.size
        if m <= 1:
            return None, None
        num_parent = [np.sum(y == c) for c in range(self.n_classes_)]
        best_gini = 1.0 - sum((n / m) ** 2 for n in num_parent)
        best_idx, best_thr = None, None

        for idx in range(self.n_features_):
            thresholds, classes = zip(*sorted(zip(X[:, idx], y)))
            num_left = [0] * self.n_classes_
            num_right = num_parent.copy()
            for i in range(1, m):
                c = classes[i - 1]
                num_left[c] += 1
                num_right[c] -= 1

                gini_left = 1.0 - sum(
                    (num_left[x] / i) ** 2 for x in range(self.n_classes_)
                )
                gini_right = 1.0 - sum(
                    (num_right[x] / (m - i)) ** 2 for x in range(self.n_classes_)
                )
                gini = (i * gini_left + (m - i) * gini_right) / m

                if thresholds[i] == thresholds[i - 1]:
                    continue
```

```

        if gini < best_gini:
            best_gini = gini
            best_idx = idx
            best_thr = (thresholds[i] + thresholds[i - 1]) / 2
    return best_idx, best_thr

def _grow_tree(self, X, y, depth=0):
    num_samples_per_class = [np.sum(y == i) for i in range(self.n_classes_)]
    predicted_class = np.argmax(num_samples_per_class)
    node = Node(predicted_class=predicted_class)

    if depth < self.max_depth:
        idx, thr = self._best_split(X, y)
        if idx is not None:
            indices_left = X[:, idx] < thr
            X_left, y_left = X[indices_left], y[indices_left]
            X_right, y_right = X[~indices_left], y[~indices_left]
            node.feature_index = idx
            node.threshold = thr
            node.left = self._grow_tree(X_left, y_left, depth + 1)
            node.right = self._grow_tree(X_right, y_right, depth + 1)
    return node

def _predict(self, inputs):
    node = self.tree_
    while node.left:
        if inputs[node.feature_index] < node.threshold:
            node = node.left
        else:
            node = node.right
    return node.predicted_class

```

Bootstrapping

```

In [79]: def resample(X, y=None, n_samples=None, random_state=None):
    if random_state is not None:
        np.random.seed(random_state)

    n_samples = n_samples or len(X) # Use the length of X if n_samples is not s
    indices = np.random.choice(np.arange(len(X)), size=n_samples, replace=True)

    X_resampled = X[indices] # Resample the feature matrix
    if y is not None:
        y_resampled = y[indices] # Resample the target vector
        return X_resampled, y_resampled
    else:
        return X_resampled

```

Random Forest Classifier

```

In [80]: class RandomForestClassifier:
    def __init__(self, n_estimators=100, max_depth=None, sample_size=0.8):
        self.n_estimators = n_estimators
        self.max_depth = max_depth
        self.sample_size = sample_size

    def fit(self, X, y):
        self.trees = []

```

```

n_samples = int(self.sample_size * len(X)) if self.sample_size < 1 else

for _ in range(self.n_estimators):
    X_sample, y_sample = resample(X, y, n_samples=n_samples)
    tree = DecisionTreeClassifier(max_depth=self.max_depth)
    tree.fit(X_sample, y_sample)
    self.trees.append(tree)

def predict(self, X):
    tree_preds = np.array([tree.predict(X) for tree in self.trees])
    return [self._majority_vote(tree_pred) for tree_pred in tree_preds.T]

def _majority_vote(self, preds):
    vote_count = Counter(preds)
    return vote_count.most_common(1)[0][0]

```

```

In [37]: # Loadin the data
df = pd.read_csv('audio_features.csv')

```

```

In [38]: print(df.head())

```

	class	file	mfcc_1	mfcc_2	mfcc_3	\
0	up	clean_471a0925_nohash_0.wav	-539.44120	34.409560	-2.751952	
1	dog	clean_e2008f39_nohash_2.wav	-465.97638	35.120487	-43.232680	
2	stop	clean_6982fc2a_nohash_0.wav	-523.65760	28.062248	-7.387421	
3	four	clean_5ebc1cda_nohash_1.wav	-471.55438	57.740578	3.400095	
4	visual	clean_af790082_nohash_3.wav	-616.45056	24.027864	11.295362	

	mfcc_4	mfcc_5	mfcc_6	mfcc_7	mfcc_8	...	chroma_cqt_6	\
0	2.365419	-3.782099	3.173902	0.986378	0.876886	...	0.319202	
1	-10.388118	-24.109978	6.047910	-2.385078	0.540033	...	0.324543	
2	-1.413885	-15.719883	4.825247	-7.887130	-0.614132	...	0.467239	
3	4.145630	-21.726675	-8.220856	-2.365875	3.310878	...	0.690343	
4	5.342732	-5.270896	-7.247482	0.075386	-6.438763	...	0.480633	

	chroma_cqt_7	chroma_cqt_8	chroma_cqt_9	chroma_cqt_10	chroma_cqt_11	\
0	0.362963	0.480808	0.418132	0.402281	0.365248	
1	0.315692	0.388141	0.241892	0.225906	0.267078	
2	0.386619	0.365450	0.348615	0.348751	0.332943	
3	0.740724	0.623760	0.524664	0.399555	0.327990	
4	0.336053	0.281622	0.267856	0.259294	0.234649	

	chroma_cqt_12	f1	f2	f3
0	0.483461	9.722557	190.612574	431.920990
1	0.204916	0.000000	380.200208	706.498029
2	0.381025	12.955628	347.177980	695.666033
3	0.316427	0.805901	167.655704	437.181824
4	0.244176	6.443128	144.238210	420.749005

[5 rows x 83 columns]

```

In [39]: print(df.columns.tolist())
df.columns = df.columns.str.strip()

class_ = df['class']
filename = df['file']

```

```
['class', 'file', 'mfcc_1', 'mfcc_2', 'mfcc_3', 'mfcc_4', 'mfcc_5', 'mfcc_6', 'mfcc_7', 'mfcc_8', 'mfcc_9', 'mfcc_10', 'mfcc_11', 'mfcc_12', 'mfcc_13', 'mfcc_delta_1', 'mfcc_delta_2', 'mfcc_delta_3', 'mfcc_delta_4', 'mfcc_delta_5', 'mfcc_delta_6', 'mfcc_delta_7', 'mfcc_delta_8', 'mfcc_delta_9', 'mfcc_delta_10', 'mfcc_delta_11', 'mfcc_delta_12', 'mfcc_delta_13', 'mfcc_delta2_1', 'mfcc_delta2_2', 'mfcc_delta2_3', 'mfcc_delta2_4', 'mfcc_delta2_5', 'mfcc_delta2_6', 'mfcc_delta2_7', 'mfcc_delta2_8', 'mfcc_delta2_9', 'mfcc_delta2_10', 'mfcc_delta2_11', 'mfcc_delta2_12', 'mfcc_delta2_13', 'chroma_1', 'chroma_2', 'chroma_3', 'chroma_4', 'chroma_5', 'chroma_6', 'chroma_7', 'chroma_8', 'chroma_9', 'chroma_10', 'chroma_11', 'chroma_12', 'spectral_centroid', 'spectral_bandwidth', 'spectral_rolloff', 'spectral_contrast_1', 'spectral_contrast_2', 'spectral_contrast_3', 'spectral_contrast_4', 'spectral_contrast_5', 'spectral_contrast_6', 'spectral_flatness', 'rms', 'zcr', 'tempo', 'pitch', 'mel_spectrogram', 'chroma_cqt_1', 'chroma_cqt_2', 'chroma_cqt_3', 'chroma_cqt_4', 'chroma_cqt_5', 'chroma_cqt_6', 'chroma_cqt_7', 'chroma_cqt_8', 'chroma_cqt_9', 'chroma_cqt_10', 'chroma_cqt_11', 'chroma_cqt_12', 'f1', 'f2', 'f3']
```

```
In [40]: print('Before dropping')
print(df.head())
print("Columns: ", df.columns.tolist())

features = df.drop(['class', 'file'], axis=1)

print('After dropping')
print(features.head())
print("Columns: ", features.columns.tolist())
```

Before dropping

	class	file	mfcc_1	mfcc_2	mfcc_3	\
0	up	clean_471a0925_nohash_0.wav	-539.44120	34.409560	-2.751952	
1	dog	clean_e2008f39_nohash_2.wav	-465.97638	35.120487	-43.232680	
2	stop	clean_6982fc2a_nohash_0.wav	-523.65760	28.062248	-7.387421	
3	four	clean_5ebc1cda_nohash_1.wav	-471.55438	57.740578	3.400095	
4	visual	clean_af790082_nohash_3.wav	-616.45056	24.027864	11.295362	

	mfcc_4	mfcc_5	mfcc_6	mfcc_7	mfcc_8	...	chroma_cqt_6	\
0	2.365419	-3.782099	3.173902	0.986378	0.876886	...	0.319202	
1	-10.388118	-24.109978	6.047910	-2.385078	0.540033	...	0.324543	
2	-1.413885	-15.719883	4.825247	-7.887130	-0.614132	...	0.467239	
3	4.145630	-21.726675	-8.220856	-2.365875	3.310878	...	0.690343	
4	5.342732	-5.270896	-7.247482	0.075386	-6.438763	...	0.480633	

	chroma_cqt_7	chroma_cqt_8	chroma_cqt_9	chroma_cqt_10	chroma_cqt_11	\
0	0.362963	0.480808	0.418132	0.402281	0.365248	
1	0.315692	0.388141	0.241892	0.225906	0.267078	
2	0.386619	0.365450	0.348615	0.348751	0.332943	
3	0.740724	0.623760	0.524664	0.399555	0.327990	
4	0.336053	0.281622	0.267856	0.259294	0.234649	

	chroma_cqt_12	f1	f2	f3
0	0.483461	9.722557	190.612574	431.920990
1	0.204916	0.000000	380.200208	706.498029
2	0.381025	12.955628	347.177980	695.666033
3	0.316427	0.805901	167.655704	437.181824
4	0.244176	6.443128	144.238210	420.749005

[5 rows x 83 columns]

Columns: ['class', 'file', 'mfcc_1', 'mfcc_2', 'mfcc_3', 'mfcc_4', 'mfcc_5', 'mfcc_6', 'mfcc_7', 'mfcc_8', 'mfcc_9', 'mfcc_10', 'mfcc_11', 'mfcc_12', 'mfcc_13', 'mfcc_delta_1', 'mfcc_delta_2', 'mfcc_delta_3', 'mfcc_delta_4', 'mfcc_delta_5', 'mfcc_delta_6', 'mfcc_delta_7', 'mfcc_delta_8', 'mfcc_delta_9', 'mfcc_delta_10', 'mfcc_delta_11', 'mfcc_delta_12', 'mfcc_delta_13', 'mfcc_delta2_1', 'mfcc_delta2_2', 'mfcc_delta2_3', 'mfcc_delta2_4', 'mfcc_delta2_5', 'mfcc_delta2_6', 'mfcc_delta2_7', 'mfcc_delta2_8', 'mfcc_delta2_9', 'mfcc_delta2_10', 'mfcc_delta2_11', 'mfcc_delta2_12', 'mfcc_delta2_13', 'chroma_1', 'chroma_2', 'chroma_3', 'chroma_4', 'chroma_5', 'chroma_6', 'chroma_7', 'chroma_8', 'chroma_9', 'chroma_10', 'chroma_11', 'chroma_12', 'spectral_centroid', 'spectral_bandwidth', 'spectral_rolloff', 'spectral_contrast_1', 'spectral_contrast_2', 'spectral_contrast_3', 'spectral_contrast_4', 'spectral_contrast_5', 'spectral_contrast_6', 'spectral_flatness', 'rms', 'zcr', 'tempo', 'pitch', 'mel_spectrogram', 'chroma_cqt_1', 'chroma_cqt_2', 'chroma_cqt_3', 'chroma_cqt_4', 'chroma_cqt_5', 'chroma_cqt_6', 'chroma_cqt_7', 'chroma_cqt_8', 'chroma_cqt_9', 'chroma_cqt_10', 'chroma_cqt_11', 'chroma_cqt_12', 'f1', 'f2', 'f3']

After dropping

	mfcc_1	mfcc_2	mfcc_3	mfcc_4	mfcc_5	mfcc_6	mfcc_7	\
0	-539.44120	34.409560	-2.751952	2.365419	-3.782099	3.173902	0.986378	
1	-465.97638	35.120487	-43.232680	-10.388118	-24.109978	6.047910	-2.385078	
2	-523.65760	28.062248	-7.387421	-1.413885	-15.719883	4.825247	-7.887130	
3	-471.55438	57.740578	3.400095	4.145630	-21.726675	-8.220856	-2.365875	
4	-616.45056	24.027864	11.295362	5.342732	-5.270896	-7.247482	0.075386	

	mfcc_8	mfcc_9	mfcc_10	...	chroma_cqt_6	chroma_cqt_7	\
0	0.876886	1.371370	1.241352	...	0.319202	0.362963	
1	0.540033	-2.484227	3.905370	...	0.324543	0.315692	
2	-0.614132	-7.394081	11.772099	...	0.467239	0.386619	
3	3.310878	-0.327795	1.879581	...	0.690343	0.740724	
4	-6.438763	-7.817415	0.397562	...	0.480633	0.336053	

	chroma_cqt_8	chroma_cqt_9	chroma_cqt_10	chroma_cqt_11	chroma_cqt_12 \
0	0.480808	0.418132	0.402281	0.365248	0.483461
1	0.388141	0.241892	0.225906	0.267078	0.204916
2	0.365450	0.348615	0.348751	0.332943	0.381025
3	0.623760	0.524664	0.399555	0.327990	0.316427
4	0.281622	0.267856	0.259294	0.234649	0.244176

	f1	f2	f3
0	9.722557	190.612574	431.920990
1	0.000000	380.200208	706.498029
2	12.955628	347.177980	695.666033
3	0.805901	167.655704	437.181824
4	6.443128	144.238210	420.749005

[5 rows x 81 columns]

Columns: ['mfcc_1', 'mfcc_2', 'mfcc_3', 'mfcc_4', 'mfcc_5', 'mfcc_6', 'mfcc_7', 'mfcc_8', 'mfcc_9', 'mfcc_10', 'mfcc_11', 'mfcc_12', 'mfcc_13', 'mfcc_delta_1', 'mfcc_delta_2', 'mfcc_delta_3', 'mfcc_delta_4', 'mfcc_delta_5', 'mfcc_delta_6', 'mfcc_delta_7', 'mfcc_delta_8', 'mfcc_delta_9', 'mfcc_delta_10', 'mfcc_delta_11', 'mfcc_delta_12', 'mfcc_delta_13', 'mfcc_delta2_1', 'mfcc_delta2_2', 'mfcc_delta2_3', 'mfcc_delta2_4', 'mfcc_delta2_5', 'mfcc_delta2_6', 'mfcc_delta2_7', 'mfcc_delta2_8', 'mfcc_delta2_9', 'mfcc_delta2_10', 'mfcc_delta2_11', 'mfcc_delta2_12', 'mfcc_delta2_13', 'chroma_1', 'chroma_2', 'chroma_3', 'chroma_4', 'chroma_5', 'chroma_6', 'chroma_7', 'chroma_8', 'chroma_9', 'chroma_10', 'chroma_11', 'chroma_12', 'spectral_centroid', 'spectral_bandwidth', 'spectral_rolloff', 'spectral_contrast_1', 'spectral_contrast_2', 'spectral_contrast_3', 'spectral_contrast_4', 'spectral_contrast_5', 'spectral_contrast_6', 'spectral_flatness', 'rms', 'zcr', 'tempo', 'pitch', 'mel_spectrogram', 'chroma_cqt_1', 'chroma_cqt_2', 'chroma_cqt_3', 'chroma_cqt_4', 'chroma_cqt_5', 'chroma_cqt_6', 'chroma_cqt_7', 'chroma_cqt_8', 'chroma_cqt_9', 'chroma_cqt_10', 'chroma_cqt_11', 'chroma_cqt_12', 'f1', 'f2', 'f3']

```
In [41]: features = features.fillna(features.mean())

print("Features Dataframe:")
print(features.head())
print("Class:")
print(class_)

class_ = class_.astype('category')

# check for missing values
print("Missing values in features: ", features.isnull().sum().sum())
print("Missing values in class: ", class_.isnull().sum())
```

Features Dataframe:

	mfcc_1	mfcc_2	mfcc_3	mfcc_4	mfcc_5	mfcc_6	mfcc_7	\
0	-539.44120	34.409560	-2.751952	2.365419	-3.782099	3.173902	0.986378	
1	-465.97638	35.120487	-43.232680	-10.388118	-24.109978	6.047910	-2.385078	
2	-523.65760	28.062248	-7.387421	-1.413885	-15.719883	4.825247	-7.887130	
3	-471.55438	57.740578	3.400095	4.145630	-21.726675	-8.220856	-2.365875	
4	-616.45056	24.027864	11.295362	5.342732	-5.270896	-7.247482	0.075386	

	mfcc_8	mfcc_9	mfcc_10	...	chroma_cqt_6	chroma_cqt_7	\
0	0.876886	1.371370	1.241352	...	0.319202	0.362963	
1	0.540033	-2.484227	3.905370	...	0.324543	0.315692	
2	-0.614132	-7.394081	11.772099	...	0.467239	0.386619	
3	3.310878	-0.327795	1.879581	...	0.690343	0.740724	
4	-6.438763	-7.817415	0.397562	...	0.480633	0.336053	

	chroma_cqt_8	chroma_cqt_9	chroma_cqt_10	chroma_cqt_11	chroma_cqt_12	\
0	0.480808	0.418132	0.402281	0.365248	0.483461	
1	0.388141	0.241892	0.225906	0.267078	0.204916	
2	0.365450	0.348615	0.348751	0.332943	0.381025	
3	0.623760	0.524664	0.399555	0.327990	0.316427	
4	0.281622	0.267856	0.259294	0.234649	0.244176	

	f1	f2	f3
0	9.722557	190.612574	431.920990
1	0.000000	380.200208	706.498029
2	12.955628	347.177980	695.666033
3	0.805901	167.655704	437.181824
4	6.443128	144.238210	420.749005

[5 rows x 81 columns]

Class:

0	up
1	dog
2	stop
3	four
4	visual

...

51762	five
51763	up
51764	six
51765	backward
51766	four

Name: class, Length: 51767, dtype: object

Missing values in features: 0


Missing values in class: 0

```
In [42]: def label_encode(class_):
class_labels = class_.cat.codes
return class_labels

X = features.values

y = label_encode(class_)
```

```
In [43]: # correlation between features
correlation = features.corr()
plt.figure(figsize=(28, 20))
sns.heatmap(correlation, cmap = 'coolwarm')
plt.show()
```



The above correlatuon matrix shows that although our features are dependent on each other and thus naive bayes might not be our best option, however the features are not extremely dependent on each other and thus we can use a decision tree classifier and random forest classifier to classify the data.

Random Forest Classifier will work the best as it will take the average of all the decision trees and thus will be able to classify the data better avoiding overfitting.

Train Test Split

```
In [44]: def train_test_split(X, y, test_size=0.2, random_state=42):
np.random.seed(random_state)
shuffle_indices = np.random.permutation(len(X))
test_set_size = int(len(X) * test_size)
test_indices = shuffle_indices[:test_set_size]
train_indices = shuffle_indices[test_set_size:]
return X[train_indices], X[test_indices], y[train_indices], y[test_indices]

X_train, X_test, y_train, y_test = train_test_split(X, y)
```

```
In [45]: X_train.shape
```

```
Out[45]: (41414, 81)
```

```
In [46]: y_train.shape
```

```
Out[46]: (41414,)
```



```
In [47]: X_test.shape
```

```
Out[47]: (10353, 81)
```

```
In [48]: y_test.shape
```

```
Out[48]: (10353,)
```

```
In [58]: # Decision Tree Classifier
```

```
dt_model = DecisionTreeClassifier(max_depth=50)  
dt_model.fit(X_train, y_train)
```

```
y_pred = dt_model.predict(X_test)
```

```
In [62]: def accuracy(y_true, y_pred):  
         return np.sum(y_true == y_pred) / len(y_true)  
  
         print("Train Accuracy: ", accuracy(y_train, dt_model.predict(X_train)))  
         print("Test Accuracy: ", accuracy(y_test, y_pred))
```

```
Train Accuracy:  0.9999758535760854
```

```
Test Accuracy:  0.31304935767410413
```

```
In [ ]:
```