

```
In [99]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import umap
from sklearn.ensemble import RandomForestClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.preprocessing import StandardScaler
from sklearn.decomposition import PCA
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, classification_report
from sklearn.preprocessing import LabelEncoder
from sklearn.ensemble import StackingClassifier
from sklearn.linear_model import Perceptron
import pickle
import xgboost as xgb
```

Creating the dataset from the csv file for model implementation

```
In [122... df = pd.read_csv('image_features.csv')
```

```
In [123... print(df.head())
# print(df.columns)
```

	image	label	0	1	2	3	\
0	Image_1.jpg	sitting	90316.0	100405.0	102866.0	101792.0	
1	Image_2.jpg	using_laptop	27590.0	28243.0	29716.0	31651.0	
2	Image_3.jpg	hugging	41103.0	44295.0	45675.0	46283.0	
3	Image_4.jpg	sleeping	51811.0	50335.0	48994.0	51602.0	
4	Image_5.jpg	using_laptop	103939.0	111915.0	106367.0	110161.0	

	4	5	6	7	...	40	41	\
0	96516.0	98304.0	96950.0	75810.0	...	0.225891	130.728963	
1	29811.0	28081.0	26216.0	28266.0	...	0.275085	109.515815	
2	47759.0	45378.0	45144.0	42223.0	...	0.235718	105.140167	
3	50489.0	50175.0	56211.0	50671.0	...	0.160645	113.777331	
4	107363.0	104932.0	106317.0	103129.0	...	0.146301	136.100860	

	42	43	44	45	46	47	\
0	132.117393	133.384722	133.184732	130.621652	128.982899	130.538279	
1	120.993655	132.245678	136.200571	137.259530	132.195489	130.230154	
2	108.586859	120.555760	133.141853	134.412390	123.543442	127.917122	
3	110.377107	121.465424	138.285705	145.717525	141.671228	134.640820	
4	135.827105	140.418473	138.276176	125.030741	123.779137	121.931994	

	48	49
0	127.256590	0.062440
1	126.812527	0.029015
2	130.942405	0.293416
3	133.754184	0.657319
4	140.837729	0.499537

[5 rows x 52 columns]

```
In [124... print(df.columns.tolist())
df.columns = df.columns.str.strip()
```

```
labels = df['label']  
image_names = df['image']
```

```
['image', 'label', '0', '1', '2', '3', '4', '5', '6', '7', '8', '9', '10', '11',  
'12', '13', '14', '15', '16', '17', '18', '19', '20', '21', '22', '23', '24', '2  
5', '26', '27', '28', '29', '30', '31', '32', '33', '34', '35', '36', '37', '38',  
'39', '40', '41', '42', '43', '44', '45', '46', '47', '48', '49']
```

In [125...

```
print("Before dropping:")  
print(df.head())  
print("Columns:", df.columns.tolist())  
  
features = df.drop(['image', 'label'], axis=1, errors='ignore')  
  
print("After dropping:")  
print(features.head())  
print("Columns:", features.columns.tolist())
```

Before dropping:

	image	label	0	1	2	3	\
0	Image_1.jpg	sitting	90316.0	100405.0	102866.0	101792.0	
1	Image_2.jpg	using_laptop	27590.0	28243.0	29716.0	31651.0	
2	Image_3.jpg	hugging	41103.0	44295.0	45675.0	46283.0	
3	Image_4.jpg	sleeping	51811.0	50335.0	48994.0	51602.0	
4	Image_5.jpg	using_laptop	103939.0	111915.0	106367.0	110161.0	

	4	5	6	7	...	40	41	\
0	96516.0	98304.0	96950.0	75810.0	...	0.225891	130.728963	
1	29811.0	28081.0	26216.0	28266.0	...	0.275085	109.515815	
2	47759.0	45378.0	45144.0	42223.0	...	0.235718	105.140167	
3	50489.0	50175.0	56211.0	50671.0	...	0.160645	113.777331	
4	107363.0	104932.0	106317.0	103129.0	...	0.146301	136.100860	

	42	43	44	45	46	47	\
0	132.117393	133.384722	133.184732	130.621652	128.982899	130.538279	
1	120.993655	132.245678	136.200571	137.259530	132.195489	130.230154	
2	108.586859	120.555760	133.141853	134.412390	123.543442	127.917122	
3	110.377107	121.465424	138.285705	145.717525	141.671228	134.640820	
4	135.827105	140.418473	138.276176	125.030741	123.779137	121.931994	

	48	49
0	127.256590	0.062440
1	126.812527	0.029015
2	130.942405	0.293416
3	133.754184	0.657319
4	140.837729	0.499537

[5 rows x 52 columns]

Columns: ['image', 'label', '0', '1', '2', '3', '4', '5', '6', '7', '8', '9', '10', '11', '12', '13', '14', '15', '16', '17', '18', '19', '20', '21', '22', '23', '24', '25', '26', '27', '28', '29', '30', '31', '32', '33', '34', '35', '36', '37', '38', '39', '40', '41', '42', '43', '44', '45', '46', '47', '48', '49']

After dropping:

	0	1	2	3	4	5	6	\
0	90316.0	100405.0	102866.0	101792.0	96516.0	98304.0	96950.0	
1	27590.0	28243.0	29716.0	31651.0	29811.0	28081.0	26216.0	
2	41103.0	44295.0	45675.0	46283.0	47759.0	45378.0	45144.0	
3	51811.0	50335.0	48994.0	51602.0	50489.0	50175.0	56211.0	
4	103939.0	111915.0	106367.0	110161.0	107363.0	104932.0	106317.0	

	7	8	9	...	40	41	42	\
0	75810.0	0.180481	0.144348	...	0.225891	130.728963	132.117393	
1	28266.0	0.002014	0.013062	...	0.275085	109.515815	120.993655	
2	42223.0	0.061157	0.096558	...	0.235718	105.140167	108.586859	
3	50671.0	0.047546	0.174988	...	0.160645	113.777331	110.377107	
4	103129.0	0.048950	0.117493	...	0.146301	136.100860	135.827105	

	43	44	45	46	47	48	\
0	133.384722	133.184732	130.621652	128.982899	130.538279	127.256590	
1	132.245678	136.200571	137.259530	132.195489	130.230154	126.812527	
2	120.555760	133.141853	134.412390	123.543442	127.917122	130.942405	
3	121.465424	138.285705	145.717525	141.671228	134.640820	133.754184	
4	140.418473	138.276176	125.030741	123.779137	121.931994	140.837729	

	49
0	0.062440
1	0.029015
2	0.293416

```
3  0.657319
4  0.499537
```

```
[5 rows x 50 columns]
```

```
Columns: ['0', '1', '2', '3', '4', '5', '6', '7', '8', '9', '10', '11', '12', '13', '14', '15', '16', '17', '18', '19', '20', '21', '22', '23', '24', '25', '26', '27', '28', '29', '30', '31', '32', '33', '34', '35', '36', '37', '38', '39', '40', '41', '42', '43', '44', '45', '46', '47', '48', '49']
```

In [126...

```
features = features.fillna(features.mean())

labels = df['label']
image_names = df['image']

print("Features DataFrame:")
print(features.head())
print("Labels:")
print(labels)

labels = labels.astype('category')

# check for missing values
print("Missing values in features:", features.isnull().sum())
print("Missing labels:", labels.isnull().sum())
```

Features DataFrame:

	0	1	2	3	4	5	6	\
0	90316.0	100405.0	102866.0	101792.0	96516.0	98304.0	96950.0	
1	27590.0	28243.0	29716.0	31651.0	29811.0	28081.0	26216.0	
2	41103.0	44295.0	45675.0	46283.0	47759.0	45378.0	45144.0	
3	51811.0	50335.0	48994.0	51602.0	50489.0	50175.0	56211.0	
4	103939.0	111915.0	106367.0	110161.0	107363.0	104932.0	106317.0	

	7	8	9	...	40	41	42	\
0	75810.0	0.180481	0.144348	...	0.225891	130.728963	132.117393	
1	28266.0	0.002014	0.013062	...	0.275085	109.515815	120.993655	
2	42223.0	0.061157	0.096558	...	0.235718	105.140167	108.586859	
3	50671.0	0.047546	0.174988	...	0.160645	113.777331	110.377107	
4	103129.0	0.048950	0.117493	...	0.146301	136.100860	135.827105	

	43	44	45	46	47	48	\
0	133.384722	133.184732	130.621652	128.982899	130.538279	127.256590	
1	132.245678	136.200571	137.259530	132.195489	130.230154	126.812527	
2	120.555760	133.141853	134.412390	123.543442	127.917122	130.942405	
3	121.465424	138.285705	145.717525	141.671228	134.640820	133.754184	
4	140.418473	138.276176	125.030741	123.779137	121.931994	140.837729	

	49
0	0.062440
1	0.029015
2	0.293416
3	0.657319
4	0.499537

[5 rows x 50 columns]

Labels:

0	sitting
1	using_laptop
2	hugging
3	sleeping
4	using_laptop
...	
12595	sitting
12596	clapping
12597	sitting
12598	dancing
12599	listening_to_music

Name: label, Length: 12600, dtype: object

Missing values in features: 0 0

1	0
2	0
3	0
4	0
5	0
6	0
7	0
8	0
9	0
10	0
11	0
12	0
13	0
14	0
15	0
16	0

```

17  0
18  0
19  0
20  0
21  0
22  0
23  0
24  0
25  0
26  0
27  0
28  0
29  0
30  0
31  0
32  0
33  0
34  0
35  0
36  0
37  0
38  0
39  0
40  0
41  0
42  0
43  0
44  0
45  0
46  0
47  0
48  0
49  0

```

dtype: int64

Missing labels: 0

```

In [127... X = features.values

label_encoder = LabelEncoder()
y_encoded = label_encoder.fit_transform(labels)

```

### UMAP for better visualization of the data and its extracted features

```

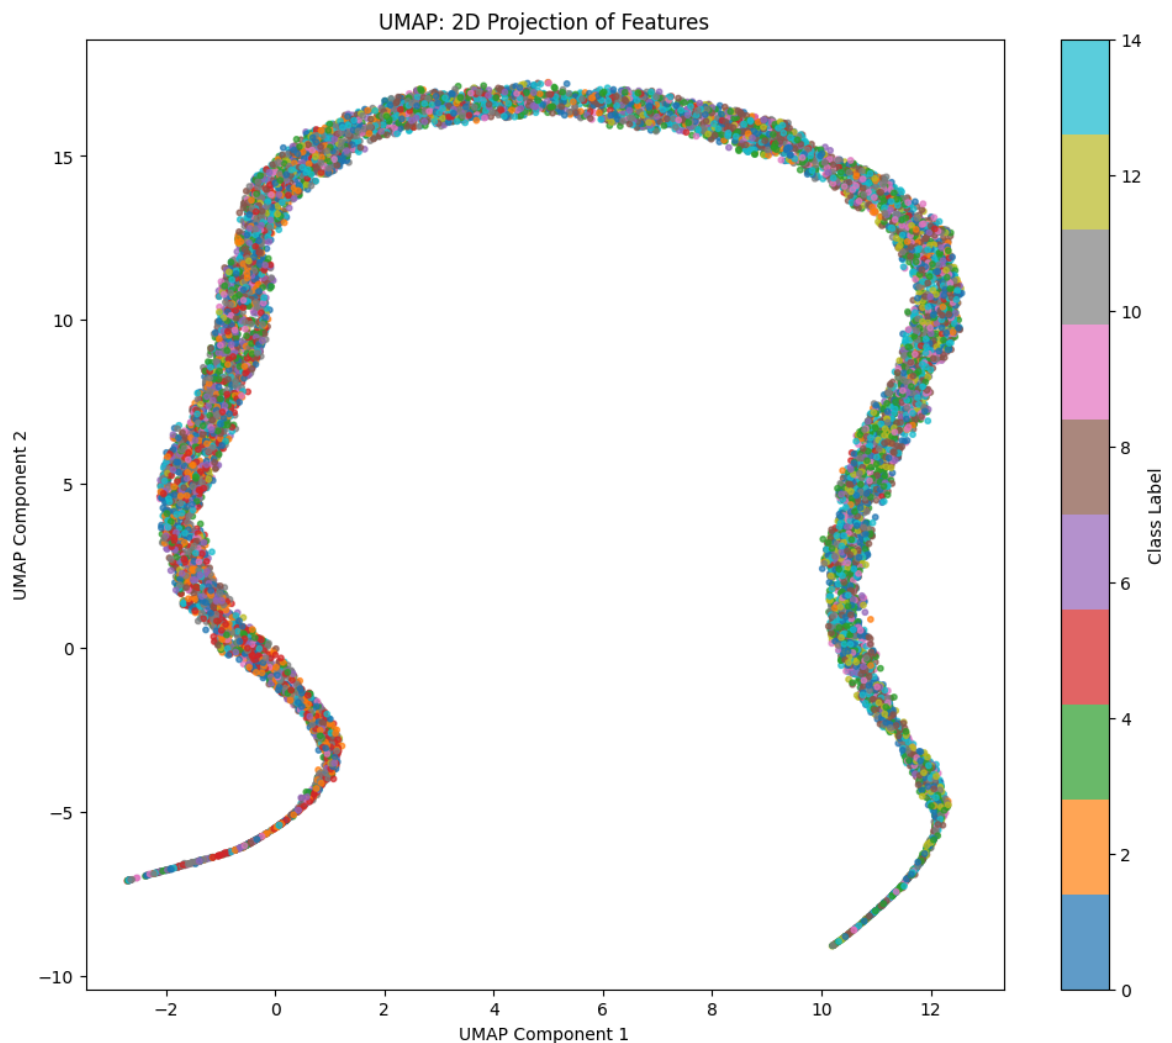
In [128... umap_model = umap.UMAP(n_components=2, random_state=42)
X_umap = umap_model.fit_transform(X)

plt.figure(figsize=(12, 10))
plt.scatter(X_umap[:, 0], X_umap[:, 1], c=y_encoded, cmap='tab10', alpha=0.7, s=
plt.title('UMAP: 2D Projection of Features')
plt.xlabel('UMAP Component 1')
plt.ylabel('UMAP Component 2')
plt.colorbar(label='Class Label')
plt.show()

```

C:\Users\Ritika\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.11\_qbz5n2kfra8p0\LocalCache\local-packages\Python311\site-packages\umap\umap.py:1945: UserWarning: n\_jobs value 1 overridden to 1 by setting random\_state. Use no seed for parallelism.

warn(f"n\_jobs value {self.n\_jobs} overridden to 1 by setting random\_state. Use no seed for parallelism.")



We can notice that our dataset is not linearly separable from the above UMAP.

### Splitting the dataset into training and testing sets 80:20

```
In [129... np.random.seed(42)
X_train, X_test, y_train, y_test = train_test_split(X, y_encoded, test_size=0.2,
```

```
In [130... X_train.shape
```

```
Out[130... (10080, 50)
```

```
In [131... X_test.shape
```

```
Out[131... (2520, 50)
```

```
In [132... y_train.shape
```

```
Out[132... (10080,)
```

```
In [133... y_test.shape
```

```
Out[133... (2520,)
```

## (3) Model Selection and Implementation

## XGBoost Classifier

XG Boost is a powerful and efficient implementation of the gradient boosting algorithm. It is an ensemble learning method that is used for classification and regression problems. It is based on the decision tree algorithm and is used to increase the accuracy of the model. It is known for its speed and performance. Theoretically, XGB should be able to perform just as well as a Random Forest, but with a lot less computational power. It is also known for its regularization techniques which help in reducing overfitting.

```
In [142... xg_model = xgb.XGBClassifier(n_estimators = 300, random_state = 42) # 275, 7
xg_model.fit(X_train, y_train)
y_pred = xg_model.predict(X_test)
```

```
In [143... xg_train_acc = accuracy_score(y_train, xg_model.predict(X_train))
print(f"Training Accuracy: {xg_train_acc* 100:.2f}%")
xg_accuracy = accuracy_score(y_test, y_pred)
print(f"Testing Accuracy: {xg_accuracy* 100:.2f}%")
print(classification_report(y_test, y_pred))
```

Training Accuracy: 100.00%

Testing Accuracy: 31.35%

	precision	recall	f1-score	support
0	0.21	0.21	0.21	173
1	0.29	0.28	0.28	160
2	0.47	0.43	0.45	186
3	0.48	0.44	0.46	180
4	0.21	0.19	0.20	152
5	0.43	0.61	0.50	151
6	0.46	0.44	0.45	186
7	0.20	0.22	0.21	151
8	0.33	0.30	0.31	179
9	0.20	0.16	0.18	176
10	0.29	0.33	0.31	155
11	0.19	0.17	0.18	163
12	0.36	0.40	0.38	162
13	0.22	0.19	0.20	183
14	0.28	0.33	0.30	163
accuracy			0.31	2520
macro avg	0.31	0.31	0.31	2520
weighted avg	0.31	0.31	0.31	2520

```
In [144... # saving the model
filename = 'xgb_model.sav'
pickle.dump(xg_model, open(filename, 'wb'))

# Loading the model
loaded_model = pickle.load(open(filename, 'rb'))
result = loaded_model.score(X_test, y_test)
print(result)
```

0.3134920634920635

## Random Forest



Theoretically, random forest would be the best model to implement for this dataset as it is a non-linear dataset and random forest is a non-linear model. Random forest is an ensemble model that uses multiple decision trees to predict the output. It is a robust model that can handle large datasets with higher dimensionality. It is also less prone to overfitting as it uses multiple decision trees to predict the output.

```
In [14]: rf_model = RandomForestClassifier(n_estimators=500, random_state=42)
rf_model.fit(X_train, y_train)
y_pred_rf = rf_model.predict(X_test)
```

```
In [15]: rf_accuracy_train = rf_model.score(X_train, y_train)
print(f"Train Accuracy for Random Forest: {rf_accuracy_train * 100:.2f}%")
rf_accuracy = accuracy_score(y_test, y_pred_rf)
print(f"Test Accuracy for Random Forest: {rf_accuracy * 100:.2f}%")
print(classification_report(y_test, y_pred_rf))
```

Train Accuracy for Random Forest: 100.00%

Test Accuracy for Random Forest: 31.75%

	precision	recall	f1-score	support
0	0.22	0.22	0.22	173
1	0.31	0.26	0.28	160
2	0.40	0.49	0.44	186
3	0.48	0.44	0.46	180
4	0.18	0.12	0.14	152
5	0.34	0.61	0.43	151
6	0.44	0.43	0.44	186
7	0.21	0.20	0.20	151
8	0.36	0.31	0.33	179
9	0.25	0.16	0.20	176
10	0.31	0.32	0.31	155
11	0.20	0.17	0.19	163
12	0.41	0.45	0.43	162
13	0.24	0.20	0.22	183
14	0.26	0.36	0.30	163
accuracy			0.32	2520
macro avg	0.31	0.32	0.31	2520
weighted avg	0.31	0.32	0.31	2520

```
In [16]: # saving the model using pickle
filename = 'random_forest_model.sav'
pickle.dump(rf_model, open(filename, 'wb'))

# Loading the model
loaded_model = pickle.load(open(filename, 'rb'))
result = loaded_model.score(X_test, y_test)
print(result)
```

0.31746031746031744

## Decision Tree

Decision tree is a simple non-linear model that can be used for classification and regression. It is a tree-like model where each node represents a feature and each branch represents a decision. It is a simple model that is easy to interpret and visualize. However,

it is prone to overfitting and may not perform well on large datasets with higher dimensionality, like the one we have.

```
In [17]: dt_model = DecisionTreeClassifier(random_state=42, max_depth=11)
dt_model.fit(X_train, y_train)
y_pred_dt = dt_model.predict(X_test)

In [18]: dt_accuracy_train = dt_model.score(X_train, y_train)
print(f"Train Accuracy for Decision Tree: {dt_accuracy_train * 100:.2f}%")
dt_accuracy = accuracy_score(y_test, y_pred_dt)
print(f"Test Accuracy for Decision Tree: {dt_accuracy * 100:.2f}%")
print(classification_report(y_test, y_pred_dt))
```

Train Accuracy for Decision Tree: 47.06%

Test Accuracy for Decision Tree: 20.32%

	precision	recall	f1-score	support
0	0.12	0.17	0.14	173
1	0.11	0.08	0.09	160
2	0.38	0.36	0.37	186
3	0.27	0.25	0.26	180
4	0.13	0.12	0.13	152
5	0.29	0.40	0.34	151
6	0.31	0.23	0.27	186
7	0.12	0.18	0.14	151
8	0.23	0.13	0.17	179
9	0.13	0.16	0.14	176
10	0.17	0.14	0.15	155
11	0.15	0.18	0.17	163
12	0.29	0.25	0.27	162
13	0.22	0.11	0.15	183
14	0.20	0.27	0.23	163
accuracy			0.20	2520
macro avg	0.21	0.20	0.20	2520
weighted avg	0.21	0.20	0.20	2520

```
In [19]: # saving the model using pickle
filename = 'decision_tree_model.sav'
pickle.dump(dt_model, open(filename, 'wb'))

# Loading the model
loaded_model = pickle.load(open(filename, 'rb'))
result = loaded_model.score(X_test, y_test)
print(result)
```

0.20317460317460317

## Scaling and PCA to reduce the dimensionality of the data and apply naive bayes

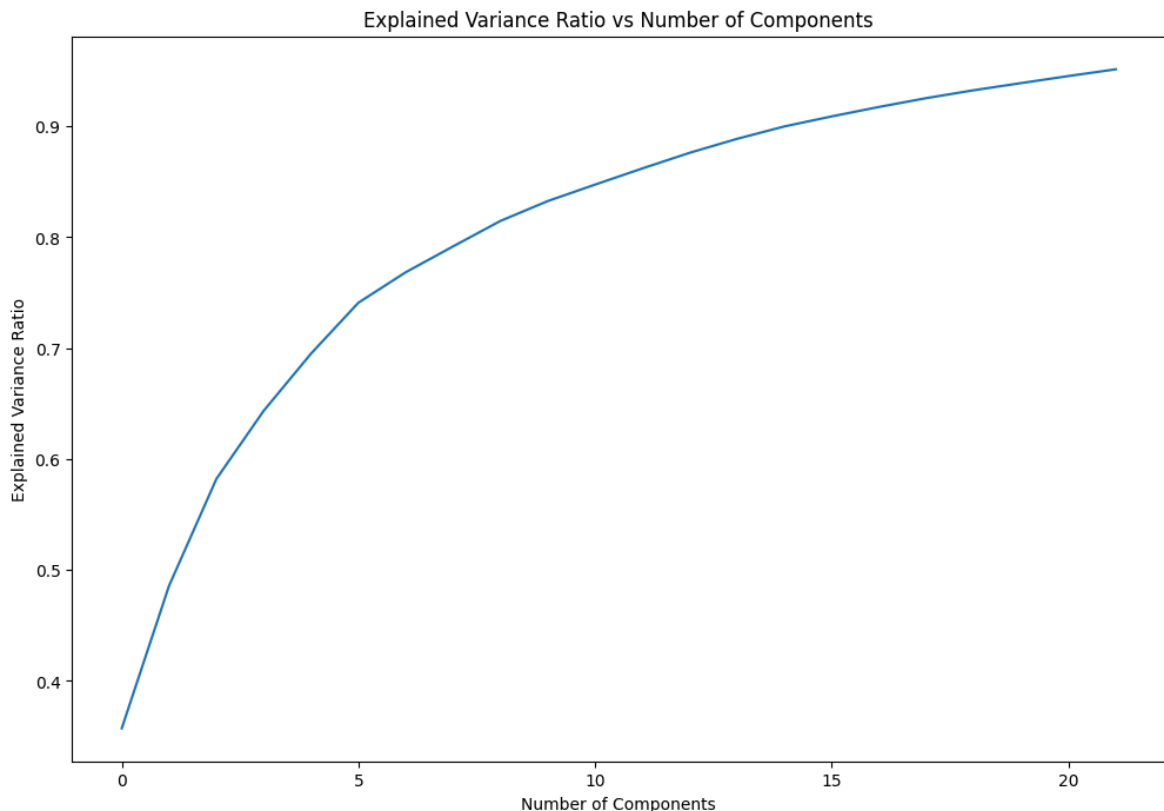
Applying PCA keeping 95% of the variance to reduce the dimensionality of the data and then applying naive bayes to classify the data.

```
In [20]: scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
```

```
X_test_scaled = scaler.transform(X_test)
```

```
In [21]: pca = PCA(n_components=0.95, random_state=42)
X_train_pca = pca.fit_transform(X_train_scaled)
X_test_pca = pca.transform(X_test_scaled)
```

```
In [22]: # plotting the explained variance ratio
plt.figure(figsize=(12, 8))
plt.plot(np.cumsum(pca.explained_variance_ratio_))
plt.xlabel('Number of Components')
plt.ylabel('Explained Variance Ratio')
plt.title('Explained Variance Ratio vs Number of Components')
plt.show()
```



## Naive Bayes

Naive Bayes is a simple probabilistic model that is based on Bayes' theorem. It is a simple model that is easy to implement and interpret. It is a good model for text classification and spam filtering. However, it assumes that the features are independent, which may not be true in real-world datasets, and is not true for our model. It is a linear model that may not perform well on non-linear datasets, like the one we have.

```
In [23]: nb_model = GaussianNB()
nb_model.fit(X_train_pca, y_train)
y_pred_nb = nb_model.predict(X_test_pca)
```

```
In [24]: nb_accuracy_train = nb_model.score(X_train_pca, y_train)
print(f"Train Accuracy for Gaussian Naive Bayes: {nb_accuracy_train * 100:.2f}%")
nb_accuracy = accuracy_score(y_test, y_pred_nb)
print(f"Test Accuracy for Gaussian Naive Bayes: {nb_accuracy * 100:.2f}%")
print(classification_report(y_test, y_pred_nb))
```

Train Accuracy for Gaussian Naive Bayes: 22.88%

Test Accuracy for Gaussian Naive Bayes: 21.19%

	precision	recall	f1-score	support
0	0.16	0.07	0.10	173
1	0.21	0.10	0.14	160
2	0.39	0.31	0.34	186
3	0.22	0.33	0.27	180
4	0.20	0.05	0.08	152
5	0.17	0.80	0.28	151
6	0.31	0.26	0.28	186
7	0.10	0.07	0.08	151
8	0.26	0.13	0.17	179
9	0.14	0.07	0.09	176
10	0.28	0.20	0.23	155
11	0.17	0.13	0.14	163
12	0.27	0.30	0.29	162
13	0.18	0.09	0.12	183
14	0.19	0.30	0.24	163
accuracy			0.21	2520
macro avg	0.22	0.21	0.19	2520
weighted avg	0.22	0.21	0.19	2520

```
In [25]: # saving the model using pickle
filename = 'gaussian_nb_model.sav'
pickle.dump(nb_model, open(filename, 'wb'))

# loading the model
loaded_model = pickle.load(open(filename, 'rb'))
result = loaded_model.score(X_test_pca, y_test)
print(result)
```

0.2119047619047619

## Stacking the models

### Random Forest + Naive Bayes + Perceptron

Stacking is an ensemble learning technique that combines multiple models to improve the performance of the final model. It is a powerful technique that can be used to improve the performance of individual models. In this case, we are stacking random forest, naive bayes, and perceptron to improve the performance of the final model. Random forest is a non-linear model that can handle large datasets with higher dimensionality. Naive bayes is a simple probabilistic model that is easy to implement and interpret. Perceptron is a simple linear model that can be used for binary classification. By combining these models, we should be able to improve the performance of the final model, atleast in comparison to both Naive Bayes and Perceptron.

```
In [26]: # stacking
estimators = [
    ('rf', RandomForestClassifier(n_estimators=500, random_state=42)),
    ('nb', GaussianNB()),
    ('perceptron', Perceptron(max_iter=800, tol=1e-3))
]
```

```
stack_model = StackingClassifier(estimators=estimators, final_estimator=RandomFo
stack_model.fit(X_train, y_train)
y_pred_stack = stack_model.predict(X_test)
```

```
In [27]: y_pred_accuracy_train = stack_model.score(X_train, y_train)
print(f"Train Accuracy for Stacking: {y_pred_accuracy_train * 100:.2f}%")
y_pred_accuracy = accuracy_score(y_test, y_pred_stack)
print(f"Test Accuracy for Stacking: {y_pred_accuracy * 100:.2f}%")
print(classification_report(y_test, y_pred_stack))
```

Train Accuracy for Stacking: 79.09%

Test Accuracy for Stacking: 28.65%

	precision	recall	f1-score	support
0	0.17	0.20	0.18	173
1	0.27	0.24	0.26	160
2	0.35	0.48	0.41	186
3	0.45	0.47	0.46	180
4	0.17	0.14	0.15	152
5	0.37	0.56	0.45	151
6	0.44	0.35	0.39	186
7	0.17	0.23	0.20	151
8	0.23	0.32	0.27	179
9	0.20	0.09	0.12	176
10	0.28	0.29	0.28	155
11	0.16	0.11	0.13	163
12	0.50	0.31	0.39	162
13	0.26	0.16	0.20	183
14	0.25	0.33	0.28	163
accuracy			0.29	2520
macro avg	0.28	0.29	0.28	2520
weighted avg	0.29	0.29	0.28	2520

```
In [28]: # saving the model using pickle
filename = 'stacking_model.sav'
pickle.dump(stack_model, open(filename, 'wb'))

# Loading the model
loaded_model = pickle.load(open(filename, 'rb'))
result = loaded_model.score(X_test, y_test)
print(result)
```

0.2865079365079365

## Naive Bayes + Random Forest

Stacking Random Forest on top of Naive Bayes should improve the performance of our model, atleast in comparison to Naive Bayes because the non-linearity of Random Forest should be able to capture the non-linear patterns in the data that Naive Bayes is not able to capture while Naive Bayes should be able to capture the linear patterns in the data that Random Forest is not able to capture, thus improving the overall performance of the model.

```
In [29]: estimators2 = [
          ('nb', GaussianNB()),
```

```

    ('rf', RandomForestClassifier(n_estimators=500, random_state=42)),
]

stack_model2 = StackingClassifier(estimators=estimators2, final_estimator=Random
stack_model2.fit(X_train, y_train)
y_pred_stack2 = stack_model2.predict(X_test)

```

```

In [30]: y_pred_accuracy_train2 = stack_model2.score(X_train, y_train)
print(f"Train Accuracy for Stacking: {y_pred_accuracy_train2 * 100:.2f}%")
y_pred_accuracy2 = accuracy_score(y_test, y_pred_stack2)
print(f"Test Accuracy for Stacking: {y_pred_accuracy2 * 100:.2f}%")
print(classification_report(y_test, y_pred_stack2))

```

Train Accuracy for Stacking: 88.12%

Test Accuracy for Stacking: 28.61%

	precision	recall	f1-score	support
0	0.20	0.20	0.20	173
1	0.28	0.23	0.25	160
2	0.36	0.44	0.40	186
3	0.44	0.47	0.45	180
4	0.15	0.13	0.14	152
5	0.39	0.57	0.46	151
6	0.43	0.38	0.40	186
7	0.16	0.19	0.17	151
8	0.24	0.27	0.25	179
9	0.20	0.10	0.14	176
10	0.26	0.29	0.27	155
11	0.16	0.15	0.16	163
12	0.43	0.31	0.36	162
13	0.23	0.17	0.20	183
14	0.27	0.36	0.31	163
accuracy				0.29
macro avg				0.28
weighted avg				0.28

```

In [31]: # saving the model using pickle
filename = 'stacking_model2.sav'
pickle.dump(stack_model2, open(filename, 'wb'))

# Loading the model
loaded_model = pickle.load(open(filename, 'rb'))
result = loaded_model.score(X_test, y_test)
print(result)

```

0.2861111111111111

## Naive Bayes + Random Forest on PCA reduced data

Since Naive Bayes performs better on lower dimensional data, we can apply PCA to reduce the dimensionality of the data and then apply Naive Bayes and Random Forest on the reduced data to improve the performance of the model.

```

In [32]: estimators3 = [
    ('nb', GaussianNB()),
    ('rf', RandomForestClassifier(n_estimators=500, random_state=42)),
]

```

```
stack_model3 = StackingClassifier(estimators=estimators3, final_estimator=Random
stack_model3.fit(X_train_pca, y_train)
y_pred_stack3 = stack_model3.predict(X_test_pca)
```

```
In [33]: y_pred_accuracy_train3 = stack_model3.score(X_train_pca, y_train)
print(f"Train Accuracy for Stacking: {y_pred_accuracy_train3 * 100:.2f}%")
y_pred_accuracy3 = accuracy_score(y_test, y_pred_stack3)
print(f"Test Accuracy for Stacking: {y_pred_accuracy3 * 100:.2f}%")
print(classification_report(y_test, y_pred_stack3))
```

Train Accuracy for Stacking: 84.74%

Test Accuracy for Stacking: 26.98%

	precision	recall	f1-score	support
0	0.17	0.14	0.15	173
1	0.27	0.24	0.26	160
2	0.35	0.39	0.37	186
3	0.38	0.42	0.39	180
4	0.13	0.11	0.12	152
5	0.34	0.58	0.43	151
6	0.36	0.37	0.36	186
7	0.14	0.15	0.14	151
8	0.27	0.31	0.29	179
9	0.16	0.09	0.11	176
10	0.25	0.32	0.28	155
11	0.19	0.13	0.15	163
12	0.39	0.34	0.36	162
13	0.25	0.16	0.20	183
14	0.22	0.31	0.25	163
accuracy			0.27	2520
macro avg	0.26	0.27	0.26	2520
weighted avg	0.26	0.27	0.26	2520

```
In [34]: # saving the model using pickle
filename = 'stacking_model3.sav'
pickle.dump(stack_model3, open(filename, 'wb'))

# Loading the model
loaded_model = pickle.load(open(filename, 'rb'))
result = loaded_model.score(X_test_pca, y_test)
print(result)
```

0.2698412698412698

Hence, we notice that our best accuracy was obtained for Random Forest Classifier. However stacking Random Forest, Naive Bayes and Perceptron gave us a quite close and decent accuracy as well. Similarly, XGBoost also gave the second best accuracy. Other models like Decision Tree and Naive Bayes did not perform well on this dataset. and needed to be stacked with other models to improve the performance.

In [ ]: